Hidden Markov Models

> **Deadline** For grade **A** you must latest have your implementations accepted by Kattis (and for the fishing derby (HMM) with the required score) by **Friday Oct 08 2021, 11:59pm sharp** and present within regular examination period (**latest** on Oct 29 2021, preferably earlier). Any assignment accepted by Kattis/presented after these deadlines will result in max. grade B.

1 Introduction to Hidden Markov Models

A Hidden Markov Model (HMM) is a powerful statistical model used in robotics, computational biology and many other disciplines. HMMs build on the idea that observations are generated by unknown or hidden states. For example, a GPS signal indicating your location is generated by your actual position. Since many processes, ranging from DNA sequences to satellite coordinates, can only be observed indirectly and under high uncertainty, HMMs are useful approximations of these systems.

HMMs belong to the class of Markov processes that describe memoryless dynamical systems. Any Markov process fulfills the Markov assumption, i.e. that the future is independent of the past given the current state at time t. In terms of HMMs, given the knowledge of the hidden state $\mathbf{X_{t-1}}$, the current hidden state $\mathbf{X_t}$ is independent of all past hidden states $\mathbf{X_\tau}$ with $\tau < t-1$. Similarly, given the current state $\mathbf{X_t}$ the current observation $\mathbf{O_t}$ is independent of all past states and observations.

In this notation, $\mathbf{X_t}$ and $\mathbf{O_t}$ are random variables which means that their value depends on a probability distribution. Realizations of these random variables, which means choosing values according to these probability distributions, are denoted by $\mathbf{x_t}$ and $\mathbf{o_t}$. Each random variable can take values from a given set of outcomes, discrete or continuous. In this assignment, we will only be dealing with discrete random variables.

As an example of that, the outcome of a coin toss can be denoted by $\mathbf{O_t}$ which can take any value in the discrete set $\{head, tail\}$. The event of observing $head$ is a realization $\mathbf{o_t}$ and has probability $P(\mathbf{O_t} = \mathbf{o_t}) = P(\mathbf{O_t} = head) = 0.5$. The probability of observing any outcome, i.e. either $head$ or $tail$, has to sum to $\sum_{\mathbf{o_t} \in \{\mathbf{head, tail}\}} P(\mathbf{O_t} = \mathbf{o_t}) = 1$.

Considering more than one time step, the notation $\mathbf{O_{1:T}} = \mathbf{o_{1:T}}$ describes an observed series of realizations of the random variable. In terms of the coin example, if we threw the coin three times and observed $head, head, tail$, then we would get $\mathbf{O_1} = head$, $\mathbf{O_2} = head$ and $\mathbf{O_3} = tail$.

Equipped with this notation, we can define the joint probability of a sequence of made observations $\mathbf{o_{1:T}}$ and all hidden states in the time interval $[1, T]$ as

$$P(\mathbf{O_{1:T}} = \mathbf{o_{1:T}}, \mathbf{X_{1:T}}) = P(\mathbf{X_1})P(\mathbf{O_1} = \mathbf{o_1}|\mathbf{X_1}) \prod_{t=2}^{T} P(\mathbf{X_t}|\mathbf{X_{t-1}})P(\mathbf{O_t} = \mathbf{o_t}|\mathbf{X_t})$$

Since the hidden states are unknown to us, they must be modeled under uncertainty. This is achieved by modeling the probability distribution over the transition from one state to another and the probability distribution over the observations given the current state. Thus, $P(\mathbf{X_t} = \mathbf{x_t}|\mathbf{X_{t-1}} = \mathbf{x_{t-1}})$ denotes

the probability of being in a specific state at time t given that the system was in a specific state at time $t-1$. $P(\mathbf{O_t} = \mathbf{o_t}|\mathbf{X_t} = \mathbf{x_t})$ is the probability of observing $\mathbf{o_t}$ given that the system is in hidden state $\mathbf{x_t}$.

At each time step $t \in [1, T]$ in the interval from 1 to T an HMM can thus be characterized by the following components:

$$\mathbf{X_t} = \mathbf{x_i}, i \in \{1, 2, ...N\} : N \text{ possible hidden states}$$

$$\mathbf{O_t} = \mathbf{o_k}, k \in \{1, 2, ...K\} : K \text{ possible observations}$$

$$\pi_i = P(\mathbf{X_1} = i) : \text{the initial probability vector } \pi \in \mathbb{R}^{1,N} \text{ with elements } \pi_i$$

$$a_{i,j} = P(\mathbf{X_{t+1}} = j|\mathbf{X_t} = i) : \text{the transition probability matrix } \mathbf{A} \in \mathbb{R}^{N,N} \text{ with elements } a_{i,j}$$

$$b_i(k) = P(\mathbf{O_t} = k|\mathbf{X_t} = i) : \text{the observation probability matrix } \mathbf{B} \in \mathbb{R}^{N,K} \text{ with elements } b_{i,k} = b_i(k)$$

As discussed above, both the hidden states and observations are random variables for which we denote the realizations by $\mathbf{X_t} = \mathbf{x_i}, i \in \{1, 2, ...N\}$ and $\mathbf{O_t} = \mathbf{o_k}, k \in \{1, 2, ...K\}$. The transition probability matrix $\mathbf{A}$ describes the probability of being in state j at time $t+1$ given that the system has been in state i at time t. Equivalently, the observation probability matrix $\mathbf{B}$ describes the probability of observing k at time t given that the system is in state i at time t. The matrix notation of the parameters simplifies the computations as the implementation of this statistical model is reduced to pure algebraic manipulations.

HMMs can be cast into different problem settings. First, if we assume that the parameter set $\lambda = (\mathbf{A}, \mathbf{B}, \pi)$ as defined above is known and that we have observed the realization $\mathbf{O_{1:T}} = \mathbf{o_{1:T}}$, we can predict the next observation. The auto-completion function for text in a mobile phone can be based on this operation - given the letters that have been observed so far, what is the most probable next letter or complete word? In this example, both the observations and hidden states are the letters of the alphabet.

Second, the probability of a given observation sequence can be estimated. Since the HMM represents the statistics of hidden states and observations, the observation of "good morning" will have a lower probability than "good morning" and an auto-correction program can detect the mistake.

Third, in order to correct the mistake, we need to decode the actual intention of the writer. Thus, we need to determine the most likely hidden state sequence that might have produced the observation sequence.

Fourth, and finally, if the parameters are not known, the set $\lambda = (\mathbf{A}, \mathbf{B}, \pi)$ needs to be learned with the help of made observations. Both the Markov assumption and the algebraic notation render this problem solvable.

1.1 GETTING STARTED

In this lab, you are required to understand and implement different HMM problems. Along with the description of your tasks, we pose a number of questions, marked by a frame in the text. You are required to answer these in written form and present them to a teaching assistant in a lab session.

For the grades E and D, this implies to implement the above-mentioned problems of HMMs, i.e. the evaluation, decoding and learning problem.

For grade C, you need to have passed the E-D level and you are furthermore required to investigate HMMs in empirical and theoretical terms.

Finally, for A and B, you need to have passed the E-D and C level. In the final assignment, you are required to apply your knowledge to an applied problem setting. We provide you with observations of the direction of swimming of different fish over time. The task is to develop a system that can determine different fish species based on their swimming pattern.

2 GRADE LEVEL E AND D

2.1 OVERVIEW

As introduced above, HMMs consist of probability distributions over hidden states and observations. In order to infer e.g. future observations or the most likely sequence of hidden states, it is essential to take any uncertainty about the system into account. This is achieved by marginalizing over states and observations at all relevant time points. For the discrete case, marginalization is achieved as follows:

In general terms, let $\mathbf{O} \sim P(\mathbf{O})$ and $\mathbf{X} \sim P(\mathbf{X})$ be random variables in a discrete space that can take the values $\mathbf{O} \in \{\mathbf{o}_1, \mathbf{o}_2, ... \mathbf{o}_K\}$ and $\mathbf{X} \in \{\mathbf{x}_1, \mathbf{x}_2, ... \mathbf{x}_N\}$ respectively. The joint probability distribution over $\mathbf{O}$ and $\mathbf{X}$ is denoted by $P(\mathbf{O}, \mathbf{X})$ and the conditional distribution of $\mathbf{O}$ given $\mathbf{X}$ by $P(\mathbf{O}|\mathbf{X})$. Then the marginal probability mass distribution of $P(\mathbf{O} = \mathbf{o}_i)$ is given by

$$P(\mathbf{O} = \mathbf{o}_i) = \sum_{j=1}^{m} P(\mathbf{O} = \mathbf{o}_i, \mathbf{X} = \mathbf{x}_j) = \sum_{j=1}^{m} P(\mathbf{O} = \mathbf{o}_i | \mathbf{X} = \mathbf{x}_j) P(\mathbf{X} = \mathbf{x}_j) \tag{2.1}$$

Here we are marginalizing over the random variable $\mathbf{X}$, i.e. in the discrete setting we sum over all its values.

As an example, imagine that we have two coins, representing the hidden states $\mathbf{X}_i$, denoted by $\mathbf{c}_1$ and $\mathbf{c}_2$, depicted in Figure 2.1. We know that coin $\mathbf{c}_1$ is biased and will show *head* with probability 0.9 and *tail* with probability 0.1, while coin $\mathbf{c}_2$ shows *head* and *tail* each with a probability of 0.5. Now assume that our observation at each time point consists of the result of a toss of one of these coins, so either $\mathbf{o}_1 = head$ or $\mathbf{o}_2 = tail$ but we do not know which of the two coins was used. In this example, the hidden states $\mathbf{X}_t$ are the identity of the coin at time t and the observations $\mathbf{O}_t$ are the observed outcome of *head* or *tail*. Let us assume that the probability of selecting any coin is 0.5, both initially and at every next time step. In order to determine the probability of observing e.g. *head* at time t we need to take into account that any of the two coins could produce this observation. Therefore, we

Figure 2.1

need to marginalize over this uncertainty. For the initial time step, $t = 1$, this gives us e.g.

$$P(\mathbf{O}_1 = head) = \sum_{i=1}^{2} P(\mathbf{O}_1 = head, \mathbf{X}_1 = \mathbf{c}_i) \tag{2.2}$$

$$= \sum_{i=1}^{2} P(\mathbf{O}_1 = head|\mathbf{X}_1 = \mathbf{c}_i)P(\mathbf{X}_1 = \mathbf{c}_i) \tag{2.3}$$

$$= P(\mathbf{O}_1 = head|\mathbf{X}_1 = \mathbf{c}_1)P(\mathbf{X}_1 = \mathbf{c}_1) + P(\mathbf{O}_1 = head|\mathbf{X}_1 = \mathbf{c}_2)P(\mathbf{X}_1 = \mathbf{c}_2) \tag{2.4}$$

$$= 0.9 * 0.5 + 0.5 * 0.5 = 0.7 \tag{2.5}$$

Question 1 This problem can be formulated in matrix form. Please specify the initial probability vector π, the transition probability matrix $\mathbf{A}$ and the observation probability matrix $\mathbf{B}$.

2.2 HMM 0 - Next observation Distribution

Imagine that we have an estimate of the current state distribution, i.e. we know which of the coins we have used at the last time step with a certain probability. In Figure 2.1, we are located at time step 7 and know the probability of using coin one or two at the previous time step. This knowledge is represented in a row vector, in which each row represents the probability to be in a certain state, here the coin we use, and the sum of all entries is 1. Given this knowledge, we can compute the probability of observing each observation, here head or tail. First, we need to multiply the transition matrix with our current estimate of states.

Question 2 What is the result of this operation?

In the following, the result must be multiplied with the observation matrix.

Question 3 What is the result of this operation?

Implement these steps in order to solve the HMM 0 problem (found in Kattis: `https://kth.kattis.com/problems/kth.ai.hmm0`).

Note: You are suggested to use Python 3, as in other assignments, but for this and the following few parts, namely HMM 0-3, you are not allowed to use the `numpy` library, as you are expected to implement all parts of the algorithm by yourself. If you would like to, feel free to implement your own matrix operations on standard python containers (lists). Keep in mind that the problems on Kattis have nothing to do with the coin example but represent arbitrary HMM models.

For testing, download the files from the homepage for the problem. You can test your python code outside Kattis by running

`pypy3 MyHMM.py < file.in`

and then make sure that the output matches that of file.ans in this case. Replace MyHMM with whatever your python file/program is called.

2.3 HMM 1 - Probability of the observation sequence

After predicting the next observation, we want to estimate what the probability of the made observation sequence $\mathbf{O}_{1:T} = [\mathbf{O}_1 = \mathbf{o}_1, \mathbf{O}_2 = \mathbf{o}_2, ..., \mathbf{O}_T = \mathbf{o}_T]$ is. For example, given our model of the coin problem, it will be extremely unlikely to observe only tails all the time.

In order to solve this problem, we can make use of the so-called forward algorithm or α-pass algorithm. This procedure iteratively estimates the probability to be in a certain state i at time t **and** having observed the observation sequence up to time t for $t \in [1, ..T]$.

We start off by computing the probability of having observed the first observation $\mathbf{o}_1$ and having been in any of the hidden states. The latter probability is provided by the initial state distribution vector π. Thus, we initialize $\alpha_1(i)$, where the subscript 1 indicates the time step and i the state, as

$$\alpha_1(i) = P(\mathbf{O}_1 = \mathbf{o}_1, \mathbf{X}_1 = \mathbf{x}_i) \qquad \text{for } i \in [1,..N] \qquad (2.6)$$

$$= P(\mathbf{O}_1 = \mathbf{o}_1 | \mathbf{X}_1 = \mathbf{x}_i) P(\mathbf{X}_1 = \mathbf{x}_i) \qquad (2.7)$$

$$= b_i(\mathbf{o}_1)\pi_i \qquad (2.8)$$

In the following steps, we need to take into account that any of the other hidden states at time $t-1$ could have proceeded the state at the current time t. In the coin example, this means that we need to take into account the probability of having used any of the coins at the last time step in order to estimate the probability of the coin that is currently used. If we would not take this probability into account, we would always guess that coin $\mathbf{c}_1$ produced a *head* observation, even if the last 6 observations have been *tail*. **Any information available to us needs to be taken into account!** Thus, at time t we need to marginalize over the probability of having been in any other state at $t-1$ and multiply this estimate with the matching observation probability as follows

$$\alpha_t(i) = P(\mathbf{O}_{1:t} = \mathbf{o}_{1:t}, \mathbf{X}_t = \mathbf{x}_i) \qquad \text{for } i \in [1,..N] \qquad (2.9)$$

$$= P(\mathbf{O}_t = \mathbf{o}_t | \mathbf{X}_t = \mathbf{x}_i, \mathbf{O}_{1:t-1}) P(\mathbf{X}_t = \mathbf{x}_i, \mathbf{O}_{1:t-1}) \qquad (2.10)$$

$$= P(\mathbf{O}_t = \mathbf{o}_t | \mathbf{X}_t = \mathbf{x}_i) P(\mathbf{X}_t = \mathbf{x}_i, \mathbf{O}_{1:t-1}) \qquad (2.11)$$

$$= P(\mathbf{O}_t = \mathbf{o}_t | \mathbf{X}_t = \mathbf{x}_i) \big[\sum_{j=1}^{N} P(\mathbf{X}_t = \mathbf{x}_i | \mathbf{X}_{t-1} = \mathbf{x}_j) P(\mathbf{X}_{t-1} = \mathbf{x}_j, \mathbf{O}_{1:t-1}) \big] \qquad (2.12)$$

$$= b_i(\mathbf{o}_t) \big[\sum_{j=1}^{N} a_{j,i} \alpha_{t-1}(j) \big] \qquad (2.13)$$

Question 4 Why is it valid to substitute $\mathbf{O}_{1:t} = \mathbf{o}_{1:t}$ with $\mathbf{O}_t = \mathbf{o}_t$ when we condition on the state $\mathbf{X}_t = \mathbf{x}_i$?

This step has to be computed iteratively for every $t \in [1,..T]$. Finally, we can compute the probability of having observed the given observation sequence $\mathbf{O}_{1:T}$. For this, we again have to marginalize over the hidden state distribution such that

$$P(\mathbf{O}_{1:T} = \mathbf{o}_{1:T}) = \sum_{j=1}^{N} P(\mathbf{O}_{1:T} = \mathbf{o}_{1:T}, \mathbf{X}_T = \mathbf{x}_j) \qquad (2.14)$$

$$= \sum_{j=1}^{N} \alpha_T(j) \qquad (2.15)$$

Implement these steps in order to solve the HMM 1 problem (found in Kattis: `https://kth.kattis.com/problems/kth.ai.hmm1`). You might want to store the α values in matrix form. Be careful to index all matrices and vectors correctly and to maintain the temporal order of summation and product as in the formulas. As before, HMM 1 is not connected to the coin example but represents arbitrary HMM problems.

2.4 HMM 2 - ESTIMATE SEQUENCE OF STATES

The forward algorithm marginalizes over the hidden state distribution. In contrast, we can also compute the most likely sequence of hidden states given the observations. Let us denote this sequence with $\{\mathbf{X}^*\}_{1:T} = (\mathbf{X}_1^*, \mathbf{X}_2^*, ... \mathbf{X}_T^*)$. Instead of assuming that any state has led to the current estimate, we only take the most likely predecessor into account. The most likely hidden state sequence can be valuable information when estimating the actual dynamics underlying our noisy observations. In the coin example, you might want to know which coin has most likely been used at each time point. In Figure 2.1, such a path is indicated by the red line.

This problem is solved with the help of the Viterbi algorithm. The procedure is divided into two steps.

At first, we need to compute the probability of having observed $\mathbf{O}_{1:t} = \mathbf{o}_{1:t}$ and being in a state $\mathbf{X}_t = \mathbf{x}_i$ given the most likely preceding state $\mathbf{X}_{t-1} = \mathbf{x}_j$ for each t. This quantity is denoted by $\delta_t(i)$, defined in equations (2.16 - 2.18). We start with the initial values

$$\delta_1(i) = P(\mathbf{O}_1 = \mathbf{o}_1, \mathbf{X}_1 = \mathbf{x}_i) \qquad \text{for } i \in [1,..N] \tag{2.16}$$
$$= b_i(\mathbf{o}_1)\pi_i. \tag{2.17}$$

The subsequent steps update the δ_t as follows

$$\delta_t(i) = max_{j\in[1,..N]} \, P(\mathbf{X}_t = \mathbf{x}_i | \mathbf{X}_{t-1} = \mathbf{x}_j) P(\mathbf{O}_{1:t-1}, \mathbf{X}^*_{1:t-2}, \mathbf{X}^*_{t-1} = \mathbf{x}_j) P(\mathbf{O}_t = \mathbf{o}_t | \mathbf{X}_t = \mathbf{x}_i) \quad \text{for } i \in [1,..N] \tag{2.18}$$

$$= max_{j\in[1,..N]} \, a_{j,i}\delta_{t-1}(j)b_i(\mathbf{o}_t). \tag{2.19}$$

In order to be able to trace back the most likely sequence later on, it is convenient to store the indices of the most likely states at each step.

$$\delta_t^{idx}(i) = argmax_{j\in[1,..N]} \, a_{j,i}\delta_{t-1}(j)b_i(\mathbf{o}_t) \text{ for } i \in [1,..N] \tag{2.20}$$

So, if the algorithm determined $\delta_t(i)$ for state i at time t to have been preceded by state k, then $\delta_t^{idx}(i) = k$.

When arriving at the end of the observation sequence, T, the probability of the most likely hidden state sequence is given by

$$P(\{\mathbf{X}^*\}_{1:T}, \mathbf{O}_{1:T} = \mathbf{o}_{1:T}) = max_{j\in[1,..N]} \, \delta_T \tag{2.21}$$

In the second step, we can backtrack the sequence $\{\mathbf{X}^*\}$ by setting

$$\mathbf{X}^*_T = argmax_{j\in[1,..N]} \, \delta_T(j) \qquad \text{and} \tag{2.22}$$
$$\mathbf{X}^*_t = \delta_{t+1}^{idx}(\mathbf{X}^*_{t+1}) \tag{2.23}$$

Implement these steps in order to solve the HMM 2 problem (found in Kattis: `https://kth.kattis.com/problems/kth.ai.hmm2`). You might want to store the δ and δ^{idx} values in matrix form.

Question 5 How many values are stored in the matrices δ and δ^{idx} respectively?

2.5 HMM 3 - ESTIMATE MODEL PARAMETERS

Up to this point, we assumed that the HMM parameters were given. In a real-world scenario, this will rarely be the case. Instead, the initial distribution and transition and observation matrices must be approximated with the help of data samples. This can be done efficiently with the help of the Baum-Welch algorithm. Luckily, we have already implemented a part of this method - the α-pass algorithm in the HMM 1 problem.

Recall that the α-pass algorithm determined the joint probability of each state at time t and all observations up to t. In order to be able to determine the probability of any state at any time point given **all** observations, we need to assess the probability of a state at time t and all future observations $[t+1:T]$. For this, we can make use of the β-pass algorithm that works similar to the α-pass algorithm but traverses the observations in the opposite direction - from the last to the first. The β is defined as

$$\beta_t(i) = P(\mathbf{O}_{t+1:T} = \mathbf{o}_{t+1:T} | \mathbf{X}_t = \mathbf{x}_i) \qquad \text{for } i \in [1,..N], \tag{2.24}$$
$$\tag{2.25}$$

i.e. the probability of observing all future observations $\mathbf{O}_{t+1:T}$ given the current state $\mathbf{X}_t = \mathbf{x}_i$. Again, we need to determine β values for all time steps and all states. For this, we initialize

$$\beta_T(i) = 1 \qquad \text{for } i \in [1,..N], \tag{2.26}$$

$$\tag{2.27}$$

as we do not intend to favour any final state over the other. In the following, we iterate backward through the observation sequence and compute the β values, defined in the following equation:

$$\beta_t(i) = P(\mathbf{O}_{t+1:T} = \mathbf{o}_{t+1:T}|\mathbf{X}_t = \mathbf{x}_i) \qquad \text{for } i \in [1,..N], \tag{2.28}$$

$$= \sum_{j=1}^{N} P(\mathbf{O}_{t+2:T} = \mathbf{o}_{t+2:T}|\mathbf{X}_{t+1} = \mathbf{x}_j) P(\mathbf{O}_{t+1} = \mathbf{o}_{t+1}|\mathbf{X}_{t+1} = \mathbf{x}_j) P(\mathbf{X}_{t+1} = \mathbf{x}_j|\mathbf{X}_t = \mathbf{x}_i) \tag{2.29}$$

$$= \sum_{j=1}^{N} \beta_{t+1}(j) b_j(\mathbf{o}_{t+1}) a_{i,j}. \tag{2.30}$$

Now that we have computed both α and β, we can estimate the probability of being in state i at time t and being in state j at time $t+1$ given all made observations and the probability of being in state i at time t given all made observations. These probabilities are given by the di-gamma function

$$\gamma_t(i,j) = P(\mathbf{X}_t = \mathbf{x}_i, \mathbf{X}_{t+1} = \mathbf{x}_j|\mathbf{O}_{1:T} = \mathbf{o}_{1:T}) \tag{2.31}$$

$$= \frac{\alpha_t(i) a_{i,j} b_j(\mathbf{O_{t+1}}) \beta_{t+1}(j)}{\sum_{k=1}^{N} \alpha_T(k)} \tag{2.32}$$

and the gamma function

$$\gamma_t(i) = P(\mathbf{X}_t = \mathbf{x}_i, |\mathbf{O}_{1:T} = \mathbf{o}_{1:T}) \tag{2.33}$$

$$= \sum_{j=1}^{N} \gamma_t(i,j). \tag{2.34}$$

Question 6 Why we do we need to divide by the sum over the final α values for the di-gamma function?

Finally, we can estimate $\lambda = (\mathbf{A}, \mathbf{B}, \pi)$ by determining the expected value of the probabilities. We have the transition estimates given by

$$a_{i,j} = \frac{\sum_{t=1}^{T-1} \gamma_t(i,j)}{\sum_{t=1}^{T-1} \gamma_t(i)} \qquad \text{for } i,j \in [1,..N], \tag{2.35}$$

the observation estimates given by

$$b_j(k) = \frac{\sum_{t=1}^{T-1} \mathbf{1}(\mathbf{O}_t = k) \gamma_t(j)}{\sum_{t=1}^{T-1} \gamma_t(j)} \qquad \text{for } j \in [1,..N], k \in [1,..K], \tag{2.36}$$

where $\mathbf{1}(\mathbf{O}_t = k)$ is the indicator function that is 1, when the argument is true and 0 otherwise.

At last, the initial probabilities are given by

$$\pi_i = \gamma_1(i) \qquad \text{for } i \in [1,..N]. \tag{2.37}$$

Now, we have everything in place in order to learn the parameters of our HMM with the Baum-Welch Algorithm:

$$\text{1. Initialize } \lambda = (\mathbf{A}, \mathbf{B}, \pi) \tag{2.38}$$

$$\text{2. Compute all } \alpha_t(i), \beta_t(i), \gamma_t(i,j), \gamma_t(i) \text{ values} \tag{2.39}$$

$$\text{3. Re-estimate } \lambda = (\mathbf{A}, \mathbf{B}, \pi) \tag{2.40}$$

$$\text{4. Repeat from 2. until convergence} \tag{2.41}$$

Implement these steps in order to solve the HMM 3 problem (found in Kattis: `https://kth.kattis.com/problems/kth.ai.hmm3`). For more details, especially regarding the initialization of the parameters (!), please refer to the Stamp tutorial (A Revealing Introduction to Hidden Markov Models by Mark Stamp).

3 GRADE LEVEL C

3.1 OVERVIEW

This part extends the HMM implementation by more empirical investigations. Given your complete Baum-Welch algorithm, you are supposed to both investigate different properties of its performance and to test different parameter settings. To this end, we provide you with an HMM model which can be used for data generation. Thus, this assignment focuses on theoretical and empirical understanding of HMMs.

3.2 EMPERICAL INVESTIGATIONS

We provide you with a file that contains a data sequence (T = 1000, 10000) generated from the following HMM dynamics:

$$\mathbf{A} = \begin{pmatrix} 0.7 & 0.05 & 0.25 \\ 0.1 & 0.8 & 0.1 \\ 0.2 & 0.3 & 0.5 \end{pmatrix} \qquad \mathbf{B} = \begin{pmatrix} 0.7 & 0.2 & 0.1 & 0 \\ 0.1 & 0.4 & 0.3 & 0.2 \\ 0 & 0.1 & 0.2 & 0.7 \end{pmatrix} \qquad \pi = \begin{pmatrix} 1 & 0 & 0 \end{pmatrix}$$

(3.1)

Your task is to train an HMM on a varying number of these observations with the help of your Baum-Welch implementation.

Question 7 Train an HMM with the same parameter dimensions as above, i.e. **A** should be a 3 times 3 matrix, etc. Initialize your algorithm with the following matrices:

$$\mathbf{A} = \begin{pmatrix} 0.54 & 0.26 & 0.20 \\ 0.19 & 0.53 & 0.28 \\ 0.22 & 0.18 & 0.6 \end{pmatrix} \qquad \mathbf{B} = \begin{pmatrix} 0.5 & 0.2 & 0.11 & 0.19 \\ 0.22 & 0.28 & 0.23 & 0.27 \\ 0.19 & 0.21 & 0.15 & 0.45 \end{pmatrix} \qquad \pi = \begin{pmatrix} 0.3 & 0.2 & 0.5 \end{pmatrix}$$

Does the algorithm converge? How many observations do you need for the algorithm to converge? How can you define convergence?

Question 8 Train an HMM with the same parameter dimensions as above, i.e. A is a 3x3 matrix etc. The initialization is left up to you.

How close do you get to the parameters above, i.e. how close do you get to the generating parameters in Eq. 3.1? What is the problem when it comes to estimating the distance between these matrices? How can you solve these issues?

Question 9 Train an HMM with different numbers of hidden states.

What happens if you use more or less than 3 hidden states? Why?

Are three hidden states and four observations the best choice? If not, why? How can you determine the optimal setting? How does this depend on the amount of data you have?

Question 10 Initialize your Baum-Welch algorithm with a uniform distribution. How does this affect the learning?

Initialize your Baum-Welch algorithm with a diagonal **A** matrix and $\pi = [0, 0, 1]$. How does this affect the learning?

Initialize your Baum-Welch algorithm with a matrices that are close to the solution. How does this affect the learning?

4 GRADE LEVEL B AND A

4.1 OVERVIEW

In this part of the assignment you will implement an agent that can classify fish. The aim is to get a deep understanding of Hidden Markov Models and for you to show that you can make use of them to solve a less structured problem. You will practice identifying the number of states, observations, etc. You will have to learn the HMM model and use it for classification.

In this part of fishing derby, you play a single-player game and observe the movement of different fish swimming in the water. There are 7 fish species, and each species has its own swimming patterns and behave differently, unknown to the player. You have to observe the swimming patterns of the fish and identify their types.

Figure 4.1 shows how the game is visualized in the environment and depicts how some fish have already been guessed correctly (full color), incorrectly (transparent) or have not been guessed yet (dark silhouettes).


Figure 4.1: Fishing derby

4.2 FISHING DERBY - GENERAL INFORMATION

This assignment requires a thorough understanding of HMMs. Therefore it is **highly recommended** to have a look at the Stamp tutorial (A Revealing Introduction to Hidden Markov Models by Mark Stamp)! It will give you more details about numerically stable computations and other important factors for your implementation.

You find valuable information regarding the dynamics of the game on Kattis: `https://kth.kattis.com/problems/kth.ai.hmm`.

PROVIDED CODE

A basic code skeleton is provided for you in Python 3. The program is fully functional as it is provided, but the program never makes any guesses. The code can be downloaded from the attachments on the Kattis page of the game, or can be downloaded directly by the following link: `https://kth.kattis.com/problems/kth.ai.hmm/file/statement/attachments/fishingderby_hmm.zip`.

The fish and the environments provided for testing are different from the ones that are used in Kattis.

You should modify the PlayerControllerHMM class in the file player.py. You may also create any number of new classes and files. The files included in the skeleton may be modified locally, but keep in mind that they will be overwritten on Kattis (except for player.py).

Your interface with the program environment is the PlayerControllerHMM class. **Avoid using stdin and stdout.** (Instead, use stderr for debugging.)

There is one test environment available to you in a sequences.json file. To test your code, you can simply run the command

```
python3 main.py < sequences.json
```

The Kattis server uses Python 3 with PyPy compiler, which unfortunately is not supported by the provided skeleton (due to the graphical user interface libraries). Although the numpy library is included in this part of the assignment on Kattis, it would be important to keep in mind that code written on python lists would actually run faster than a code that uses numpy.

For more details, check out the problem description on Kattis webpage.

GRADING DETAILS AND TASKS

You are required to have a **Kattis score of at least 250**. However, keep in mind that the goal of this exercise is not to optimize your code toward Kattis, but instead that you get an **understanding of the theory behind HMMs**. Therefore, you also have to be able to reason about your solution and answer theoretical questions regarding HMMs for obtaining your grade.